

Brain Tumor MRI Analysis System

AI-Driven Classification · Interpretability · Clinical Decision Support

~3,500

MRI Scans

4

Tumor Classes

94–96%

Ensemble Accuracy

3

DL Models

Presentation Agenda

01

Problem Statement & Objectives

Real-world clinical challenges and mission

02

Dataset Overview

~3,500 MRI scans, 4 tumor classes, imbalance analysis

03

System Architecture

Three-model ensemble pipeline

04

Data Processing & Training

Augmentation, focal loss, AdamW optimization

05

Evaluation & Interpretability

GradCAM, attention maps, uncertainty estimation

06

Performance Results

Per-class F1, AUC-ROC, ensemble metrics

07

Clinical Workflow & Compliance

Radiologist integration, HIPAA, regulatory notes

08

Conclusion

Summary and deployment readiness

Problem Statement & Objectives

Mission

Build a production-ready AI/ML system for robust brain tumor classification from MRI images—going beyond basic classification to address real-world clinical challenges.

Class Imbalance

2.81× ratio — handled via weighted/focal loss

Inter-patient Variability

Different MRI machines & acquisition protocols

Noise & Quality

Real-world MRI artifacts and resolution differences

Limited Annotations

Semi-supervised and transfer learning strategies

Advanced Objectives



Tumor Localization via Attention Maps



Uncertainty Estimation & Confidence



GradCAM Interpretability



Ensemble Robustness




Semi-supervised / Transfer Learning



Clinical Workflow Integration

Dataset Overview



Class Distribution





Image Properties

Resolution:	512×512 to 1024×1024 px
Modality:	T1/T2 Weighted MRI
Format:	Grayscale (converted RGB)
Protocol:	Multiple acquisition types
Noise:	Real-world artifacts present
Source:	Multi-center clinical data

System Architecture



Baseline

ResNet50

- ▶ 50 residual layers + skip connections
- ▶ ImageNet pretrained weights
- ▶ ~25M parameters
- ▶ Accuracy: 92–94%

Efficient

EfficientNet-B3

- ▶ Compound scaling (depth/width/res)
- ▶ Depthwise separable convolutions
- ▶ 50% fewer params vs ResNet50
- ▶ Accuracy: 93–95%

Interpretable

ResNet50 + Attention

- ▶ Spatial attention mechanism
- ▶ Highlights tumor regions
- ▶ Enhanced GradCAM output
- ▶ Accuracy: 92–94%

Data Processing & Training Strategy

Data Processing & Augmentation

Preprocessing



Normalization

Zero-mean, unit-variance scaling



CLAHE Enhancement

Adaptive histogram equalization for contrast



Stratified Split

80% train / 10% val / 10% test

Augmentation Techniques (8+)



Random Rotation ($\pm 30^\circ$)



Horizontal / Vertical Flip



Affine Transforms



Color Jitter



Elastic Deformation



Random Crop / Zoom

Training Strategy

Loss Functions

- Cross-Entropy with class weights (inverse freq)
- Focal Loss — focus on hard/misclassified examples

Optimization

- AdamW optimizer with weight decay (L2)
- ReduceLROnPlateau scheduling
- Gradient clipping (max norm = 1.0)

Regularization

- Dropout layers ($p=0.3-0.5$)
- L2 weight decay in AdamW
- Early stopping on validation loss

Evaluation & Interpretability

Evaluation Metrics

- Accuracy:** Overall correct classification rate
- Precision & Recall:** Per-class performance analysis
- F1-Score:** Harmonic mean of precision/recall
- AUC-ROC:** Multi-class one-vs-rest curves
- Confusion Matrix:** Full cross-class error analysis

Interpretability Techniques

GradCAM

Gradient-based class activation maps show which regions triggered classification

Spatial Attention

Learned attention weights focus on tumor-bearing areas in the scan

Feature Visualization

t-SNE / UMAP of learned embeddings for cluster analysis

Uncertainty Estimation

Entropy

Shannon entropy over class probabilities. High entropy = uncertain prediction needing expert review.


Margin Score

Gap between top-2 probabilities. Narrow margin signals ambiguous cases between classes.


Confidence Intervals

Bootstrap or ensemble-based bounds on prediction reliability for clinical decision thresholds.

Performance Results



Model Accuracy Comparison



Clinical Workflow Integration

1

Patient MRI Scan

Acquire T1/T2 images

2

Preprocessing

Normalize + CLAHE enhance

3

Model Inference

Prediction + Confidence

4

Confidence Check

Threshold evaluation

5

Radiologist Review

GradCAM visualization

6

Clinical Decision

Model + expert judgment

Advanced Concepts Implemented

Class Imbalance Handling

- ▶ Weighted Cross-Entropy Loss (inverse class frequency)
- ▶ Focal Loss — down-weights easy examples
- ▶ Balanced batch sampling strategies
- ▶ Stratified K-fold cross-validation

Transfer Learning

- ▶ ImageNet pretrained ResNet50 & EfficientNet-B3
- ▶ Fine-tuned on MRI domain data
- ▶ Reduces labeled data requirements
- ▶ Strong generalization from limited samples

Spatial Attention Mechanism

- ▶ Learns to focus on tumor-relevant regions
- ▶ Channel + spatial attention gates
- ▶ Better GradCAM visualizations
- ▶ Improved clinical interpretability

Ensemble & Robustness

- ▶ Averages softmax probabilities across 3 models
- ▶ Reduces variance and model-specific biases
- ▶ Confidence bounds from prediction spread
- ▶ More stable on edge-case/noisy images

Project Structure & Technology Stack

Project Files

- ▶ **00_run_all.py**
Master pipeline script
- ▶ **01_data_exploration.py**
Dataset analysis & stats
- ▶ **02_train_models.py**
Model training workflow
- ▶ **03_evaluate_models.py**
Metrics & GradCAM
- ▶ **04_clinical_report.py**
Clinical recommendations
- ▶ **utils_helpers.py**
Shared utility functions
- ▶ **requirements.txt**
Python dependencies

Technology Stack

Deep Learning

- PyTorch 2.0+
- Torchvision (pretrained)

Computer Vision

- OpenCV
- Pillow (PIL)
- NumPy

Data & Analysis

- Pandas
- Scikit-learn
- JobLib

Visualization

- Matplotlib
- Seaborn

Conclusion

Production-Ready & Clinically Relevant

Beyond Basic Classification

Localization, uncertainty, anomaly detection

Full Interpretability

GradCAM + attention maps + feature visualization

Clinical Workflow Ready

Confidence gates + radiologist integration

Class Imbalance Solved

Weighted loss + focal loss + stratified sampling

Ensemble Accuracy: 94–96%

3-model averaging for robust predictions

HIPAA & Compliance Ready

Regulatory considerations documented